

2. Bayes and Kalman Filters

2. Bayes and Kalman Filters

Bayes Filter
Formulation
Motion Modeling
Measurement Modeling
Combined Model
[Example] Discrete State Motion & Measurement Model:
Aim of Bayes Filter
Bayes Filter | Problem Formulation
Bayes Filter | Algorithm
[Bayes Filter | Algorithm Abstract]
1. **Prediction Update** : (*Total Probability*)
2. **Measurement Update** : (*Bayes Theorem*)
[Recall | Bayes Filter Theorem]
[Derivation | Proof by Induction]
Summary:
[Example 1 | Discrete Bayes Filter]
[Example 2 | Histogram Filter]
Kalman Filter
Kalman Filter Modeling Assumption
Full Model Formulation:
Belief is Gaussian:
Goal:
[KF Algorithm Abstract]:
1. Prediction Update:
2. Measurement Update:
[Example]
[Summary]
Summary:
Relation to Bayes Filter:
Relation | Problem Formulation:
Relation | Algorithm:
1. **Prediction update** (*Total Probability*):
2. **Measurement Update** (*Bayes Theorem*):
Variable Summary:
Extended Kalman Filter
EKF | Modeling Assumption:
[EKF Algorithm Abstract]
0. **Linearization with First Order Taylor Series Expansion**
1. **Prediction Update**
2. **Measurement Update**
Sample Code:
Summary
[Summary | Comparison Table]

Bayes Filter

- The Bayes Filter forms the **foundation** for all other filters in this class
 - As described in background slides, Bayes rule is the right way to incorporate new probabilistic information into an existing, **prior estimate**
 - The resulting filter definition can be implemented **directly for discrete state systems**
 - For continuous states, need additional assumptions, additional structure to **solve the update equations analytically**

Formulation

- State: x_i
 - All aspects of the vehicle and its environment that can impact the future
 - **Assume the state is complete**
- Control inputs: u_t
 - All elements of the vehicle and its environment that can be controller
- Measurements: y_t
 - All elements of the vehicle and its environment that can be sensed
- Notation:
 - Discrete time index t
 - Initial state is x_0
 - First, apply control action u_1
 - Move to state x_1
 - Then, take measurement y_1

Motion Modeling

- **Complete state:**


- At each time t , x_{t-1} is a sufficient summary of all previous inputs and measurements:
 - $p(x_t | x_{0:t-1}, y_{1:t-1}, u_{1:t}) = p(x_t | x_{t-1}, u_t)$
- Application of Conditional Independence
 - No additional information is to be had by considering previous inputs or measurements
- Referred to as the **Markov Assumption**
 - Motion model is a **Markov Chain**

Measurement Modeling

- Complete state:
 - Current state is sufficient to model all previous states, measurements and inputs:
 - $p(y_t | x_{0:t}, y_{1:t-1}, u_{1:t}) = p(y_t | x_t)$
- Again, conditional independence
- Recall, in standard LTI state space model, measurement model **may also depend on the current input**


Combined Model

- Referred to as **Hidden Markov Model (HMM)** or **Dynamic Bayes Network (DBN)**



[Example] Discrete State Motion & Measurement Model:

- Example Motion Model:
 - States: $\{NoRain, Drizzle, Steady, Downpour\}$
 - Inputs: **None**



- For discrete states, the motion model can be written in matrix form

- For each input u_t , the $n \times n$ motion model matrix is

$$p(x_t | u_t = u, x_{t-1}) = \begin{bmatrix} p(x_t = x_1 | x_{t-1} = x_1) & p(x_t = x_1 | x_{t-1} = x_2) & \dots \\ p(x_t = x_2 | x_{t-1} = x_1) & p(x_t = x_2 | x_{t-1} = x_2) & \dots \\ \vdots & \vdots & \ddots \end{bmatrix}$$

- \downarrow Each **row** — defines the probabilities of transitioning to state x_t from all possible states x_{t-1}
- \rightarrow Each **column** | defines the probabilities of transitioning to any state x_t from a specific state x_{t-1}
- Again, the columns must sum to 1: $\sum_i p_i = 1$
- Example:

$$\blacksquare p(x_t | u_t = u, x_{t-1}) = \overbrace{\begin{bmatrix} 0.8 & 0.3 & 0.05 & 0 \\ 0.1 & 0.4 & 0 & 0 \\ 0.1 & 0.3 & 0.9 & 0.5 \\ 0 & 0 & 0.05 & 0.5 \end{bmatrix}}^{x_{t-1}} \bigg\} x_t$$

- Example Measurement Model:
 - States: $\{NoRain, Drizzle, Steady, Downpour\}$
 - Measurements: $\{Dry, Light, Medium, Heavy\}$

$$\blacksquare p(x_t | u_t = u, x_{t-1}) = \overbrace{\begin{bmatrix} 0.95 & 0.1 & 0 & 0 \\ 0.05 & 0.8 & 0.15 & 0 \\ 0 & 0.1 & 0.7 & 0.1 \\ 0 & 0 & 0.15 & 0.9 \end{bmatrix}}^{x_t} \bigg\} y_t$$

\blacksquare Again, the columns must sum to 1 : $\sum_i p_i = 1$

Aim of Bayes Filter

1. To estimate the current state of the system based on all known inputs and measurements.
 - That is, to define a **belief** about the **current state** using all **available** information:
 - $\blacksquare \bar{bel}(x_t) = p(x_t | y_{1:t}, u_{1:t})$
 - Known as **belief**, state of knowledge, information state Depends on every bit of information that exists up to time t
2. Can also **define a belief prior** to **measurement** y_t
 - $\bar{bel}(x_t) = p(x_t | y_{1:t-1}, u_{1:t})$
 - Known as **prediction**, **predicted state**

Bayes Filter | Problem Formulation

- Given a **prior** for the system state:
 - $p(x_0)$
- Given **motion** and **measurement models**:
 - $\overbrace{p(x_t | x_{t-1}, u_t)}^{\text{motion}} \quad \overbrace{p(y_t | x_t)}^{\text{measurement}}$
- Given a sequence of inputs and measurements:
 - $u_{1:t} = \{u_1, \dots, u_t\}, \quad y_{1:t} = \{y_1, \dots, y_t\}$
- Estimate the current state distribution
 - (Form a **belief** about the current state):
 - $bel(x_t) = p(x_t | y_{1:t}, u_{1:t})$

Bayes Filter | Algorithm

[Bayes Filter | Algorithm Abstract]

- At each time step, t , for all possible values of the state x :

1. Prediction Update : (Total Probability)

$$\bullet \bar{bel}(x_t) = \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$$

2. Measurement Update : (Bayes Theorem)

- $bel(x_t) = \eta p(y_t | x_t) \bar{bel}(x_t)$
- η : normalizing constant
 - does not depend on the state
- Recursive estimation technique

[Recall | Bayes Filter Theorem]

- $p(a|b) = \frac{p(b|a)p(a)}{p(b)}$
- Terminology:
 - $posterior = \frac{likelihood \cdot prior}{evidence}$

[Derivation | Proof by Induction]

- Demonstrate that belief at time t can be found using
 - belief at time $t - 1$,
 - input at t
 - and measurement at t
- Initially:
 - $bel(x_0) = p(x_0)$
- At time t , from [\[Recall\] Bayes Filter Theorem](#) relates x_t, y_t

$$\begin{aligned}
 bel(x_t) &= p(x_t | y_{1:t}, u_{1:t}) = p(x_t | y_t, y_{1:t-1}, u_{1:t}) \\
 &= \frac{p(y_t | x_t, y_{1:t-1}, u_{1:t}) p(x_t | y_{1:t-1}, u_{1:t})}{p(y_t | y_{1:t-1}, u_{1:t})} \\
 &\quad \underbrace{p(y_t | x_t, y_{1:t-1}, u_{1:t})}_{p(y_t|x_t): \text{measurement model.}} \underbrace{p(x_t | y_{1:t-1}, u_{1:t})}_{\overline{bel}(x_t): \text{belief prediction.}} \underbrace{p(y_t | y_{1:t-1}, u_{1:t})^{-1}}_{\eta: \text{const. normalizer}} \\
 &= \overbrace{p(y_t | x_t, y_{1:t-1}, u_{1:t})}^{p(y_t|x_t): \text{measurement model.}} \overbrace{p(x_t | y_{1:t-1}, u_{1:t})}^{\overline{bel}(x_t): \text{belief prediction.}} \overbrace{p(y_t | y_{1:t-1}, u_{1:t})^{-1}}^{\eta: \text{const. normalizer}} \\
 &= \eta p(y_t | x_t) \overline{bel}(x_t)
 \end{aligned}$$

- Hence:
 - **Measurement Update** $\equiv bel(x_t) = \eta p(y_t | x_t) \overline{bel}(x_t)$
 - Multiplication of two vectors
 - Requires the measurement y_t to be known
- However, we now need to find the **belief prediction**:
 - Done using total probability over previous state:

$$\begin{aligned}
 \overline{bel}(x_t) &= p(x_t | y_{1:t-1}, u_{1:t}) \\
 &= \int \overbrace{p(x_t | x_{t-1}, y_{1:t-1}, u_{1:t})}^{\text{a) can incorporate motion model}} \overbrace{p(x_{t-1} | y_{1:t-1}, u_{1:t})}^{\text{b) } = \text{bel}(x_{t-1})} dx_{t-1}
 \end{aligned}$$

a) Incorporating **Motion Modeling**:

$$p(x_t | x_{t-1}, y_{1:t-1}, u_{1:t}) = p(x_t | x_{t-1}, u_t)$$

b) And we note that the control input at time t does not affect the state at time $t - 1$:

$$p(x_{t-1} | y_{1:t-1}, u_{1:t}) = p(x_{t-1} | y_{1:t-1}, u_{1:t-1}) = bel(x_{t-1})$$

- Hence:
 - **Prediction Update** $\equiv \overline{bel}(x_t) = \int p(x_t | x_{t-1}, u_t) bel(x_{t-1}) dx_{t-1}$

- END OF **proof by induction** for [Bayes Filter Algorithm](#)
 - For this step, we need the control input to define the correct motion model distribution

Summary:


- If **state, measurements, inputs** are **DISCRETE**,
 - can **directly implement Bayes Filter**
 - **Prediction update** is **summation** over **discrete** states
 - **Measurement update** is **multiplication** of **two vectors**
- Else, if they are **CONTINUOUS**
 - must define model or approximation to enable computation:
 - **KF** (Kalman Filter)
 - **Linear** motion models
 - **Linear** measurement models
 - **Additive Gaussian** disturbance and noise distributions

- **EKF / UKF** (Extended Kalman Filter / Unscented Kalman Filter)
 - **Non-linear** motion models
 - **Non-linear** measurement models
 - **Additive Gaussian** disturbance and noise distributions
- **PF** (Particle Filter)
 - **(Dis)continuous** motion models
 - **(Dis)continuous** measurement models
 - **General** disturbance and noise distributions

[Example 1 | Discrete Bayes Filter]

- **Problem:**

- Detect if a door is open/closed with a robot that can sense the door position and pull the door open



- **Formulation:**

- **State:** $door = \{open, closed\}$

- **State Prior** (uniform):

$$p(x_0) = \begin{cases} p(open) = 0.5 \\ p(closed) = 0.5 \end{cases}$$

- **Inputs:** $arm_command = \{none, pull\}$

- **Motion Model:**

- If input = *none*, do nothing:

$$p(x_t | u_t = none, x_{t-1}) \rightarrow \begin{cases} p(open_t | none, open_{t-1}) = 1 \\ p(closed_t | none, open_{t-1}) = 0 \\ p(open_t | none, closed_{t-1}) = 0 \\ p(closed_t | none, closed_{t-1}) = 1 \end{cases}$$

- If input = *pull*, pull the door open:

$$p(x_t | u_t = pull, x_{t-1}) \rightarrow \begin{cases} p(open_t | none, open_{t-1}) = 1 \\ p(closed_t | none, open_{t-1}) = 0 \\ p(open_t | none, closed_{t-1}) = 0.8 \\ p(closed_t | none, closed_{t-1}) = 0.2 \end{cases}$$

- **Measurements:** $meas = \{sense_open, sense_closed\}$

- **Measurement model** (noisy door sensor):

$$p(y|x) \rightarrow \begin{cases} p(sense_open|open) = 0.6 \\ p(sense_open|closed) = 0.2 \\ p(sense_closed|open) = 0.4 \\ p(sense_closed|closed) = 0.4 \end{cases}$$

- **Example:**

- At $t = 1$,

- input $u_1 = none$:

- **1. Prediction Update : (Total Probability):**

$$\overline{bel}(x_t) = \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$$

$$\overline{bel}(x_1) = \int p(x_1 | u_1, x_0) bel(x_0) dx_0 = \sum p(x_1 | u_1, x_0) p(x_0)$$

- Calculate belief prediction for each possible value of state:

$$\begin{aligned} \overline{bel}(open_1) &= p(open_1 | none_1, open_0) bel(open_0) + p(open_1 | none_1, closed_0) bel(closed_0) \\ &= 1 * 0.5 + 0 * 0.5 = 0.5 \end{aligned}$$

$$\begin{aligned} \overline{bel}(closed_1) &= p(closed_1 | none_1, open_0) bel(open_0) + p(closed_1 | none_1, closed_0) bel(closed_0) \\ &= 0 * 0.5 + 1 * 0.5 = 0.5 \end{aligned}$$


- measurement $y_1 = sense_open$:

- **2. Measurement Update : (Bayes Theorem):**

- $bel(x_1) = \eta p(y_1|x_1)\overline{bel}(x_1)$
- Calculate for each possible value of state:
 - $bel(open_1) = \eta p(sense_open_1|open_1)\overline{bel}(open_1)$
 $= \eta 0.6 \cdot 0.5 = 0.3\eta$
 - $bel(closed_1) = \eta p(sense_open_1|closed_1)\overline{bel}(closed_1)$
 $= \eta 0.2 \cdot 0.5 = 0.1\eta$
- Calculate **normalizer** η and solve for **posterior**
 - $\eta = \frac{1}{0.3+0.1} = 2.5$
 - $bel(open_1) = 0.75$
 - $\Rightarrow bel(closed_1) = 0.25$
- At $t = 2$,
 - With $u_2 = pull$, $y_2 = sense_open$
 - Then:
 - State propagation:
 - $\overline{bel}(open_2) = 1 \times 0.75 + 0.8 \times 0.25 = 0.95$
 - $\overline{bel}(closed_2) = 0 \times 0.75 + 0.2 \times 0.25 = 0.05$
 - Measurement Update:
 - $bel(open_2) = \eta 0.6 \cdot 0.95 = 0.983$
 - $bel(closed_2) = \eta 0.2 \cdot 0.05 = 0.017$
- In summary:
 - Uniform prior, do nothing, measure open: $bel(open_1) = 0.75$
 - Pull open, measure open: $bel(open_2) = 0.983$

[Example 2 | Histogram Filter]

- Motion of robot in a $n \times n$ grid:



- **State:**
 - Position = $\{x_{11}, x_{12}, \dots, x_{1n}, \dots, x_{nn}\}$
- **Input:**
 - Move = $\{Up, Right, Down, Left\}$
 - 40% chance the move does not happen
 - Cannot pass through outer walls
- **Measurement:** Accurate to within 3×3 grid
 - $p(y(i-1:i+1, j-1:j+1)|x(i, j)) = \begin{bmatrix} .11 & .11 & .11 \\ .11 & .12 & .11 \\ .11 & .11 & .11 \end{bmatrix}$
- **Prior over states**
 - Assume no information, uniform
 - Vector of length n^2
 - $p(x_0) = \frac{1}{n^2}$
- **Motion Model:**
 - Given a particular input and previous state, probability of moving to any other state
 - $n \times n$ state, one for each grid point
 - 4 input choices

$$\blacksquare p(x_t | x_{t-1}, u_t) \in [0, 1]^{n^2 \times n^2 \times 4}$$

▪ Code:

```

1 | mot_mod = zeros(N,N,4);
2 |
3 | for i=1:n
4 |     for j=1:n
5 |         cur = i+(j-1)*n;
6 |
7 |         % Move up
8 |         if (j > 1)
9 |             mot_mod(cur-n,cur,1) = 0.6;
10 |            mot_mod(cur,cur,1) = 0.4;
11 |         else
12 |             mot_mod(cur,cur,1) = 1;
13 |         end
14 |
15 |         % Move right
16 |         if (i < n)
17 |             mot_mod(cur+1,cur,2) = 0.6;
18 |             mot_mod(cur,cur,2) = 0.4;
19 |         else
20 |             mot_mod(cur,cur,2) = 1;
21 |         end
22 |     end
23 | end

```

◦ Measurement Model

- Given any current state, probability of a measurement
- Same number of measurements as states $p(y_t | x_t) \in [0, 1]^{n^2 \times n^2}$
- Same 3×3 matrix governs all interior points
- Boundaries cut off invalid measurements and require normalization
- Very simplistic and bloated model
 - Could replace with 2 separate states and measurements to perpendicular walls
- Code:

```

1 | %% Create the measurement model
2 |
3 | meas_mod_rel = [0.11 0.11 0.11;
4 |                0.11 0.12 0.11;
5 |                0.11 0.11 0.11];
6 |
7 | % Convert to full measurement model
8 | % p(y_t | x_t)
9 | meas_mod = zeros(N,N);
10 |
11 | % Fill in non-boundary measurements
12 | for i=2:n-1
13 |     for j=2:n-1
14 |         cur = i+(j-1)*n;
15 |         meas_mod(cur-n+[-1:1:1],cur) = meas_mod_rel(1,:);
16 |         meas_mod(cur+[-1:1:1],cur) = meas_mod_rel(2,:);
17 |         meas_mod(cur+n+[-1:1:1],cur) = meas_mod_rel(3,:);
18 |     end
19 | end

```

• Making moves:

```

1 | videoobj=VideoWriter('bayesgrid.mp4','MPEG-4');
2 | truefps = 1; videoobj.FrameRate = 10; %Anything less than 10 fps fails. open(videoobj);
3 |
4 | figure(1);clf; hold on;
5 | beliefs = reshape(bel,n,n);
6 | imagesc(beliefs);
7 | plot(pos(2),pos(1),'ro','MarkerSize',6,'LineWidth',2)
8 | colormap(bone);
9 | title('True state and beliefs')
10 | F = getframe;
11 |
12 | % Dumb hack to get desired framerate
13 | for dumb=1:floor(10/truefps)
14 |     writeVideo(videoobj, F);
15 | end

```

• Simulation Code:

```

1 | %Main Loop
2 | for t=1:T
3 |     %% Simulation
4 |     % Select motion input
5 |     u(t) = ceil(4*rand(1));
6 |     % Select a motion
7 |     thresh = rand(1);
8 |     new_x = find(cumsum(squeeze(mot_mod(:, :, u(t))))*x(:,t))>thresh,1);
9 |     % Move vehicle
10 |    x(new_x,t+1) = 1;
11 |    % Take measurement
12 |    thresh = rand(1);
13 |    new_y = find(cumsum(meas_mod(:, :, x(:,t+1)))>thresh,1);
14 |    y(new_y,t) = 1;
15 |    % Store for plotting

```

16 | ...

- Bayes Filter:

```
1 | ...
2 |
3 | %% Bayesian Estimation
4 | % Prediction update
5 | belp = squeeze(mot_mod(:,:,u(t)))*bel;
6 |
7 | % Measurement update
8 | bel = meas_mod(new_y,:)'.*belp;
9 | bel = bel/norm(bel);
10 |
11 | [pmax y_bel(t)] = max(bel);
12 |
13 | %% Plot beliefs
14 | ...
```

Kalman Filter

- Kalman convinced NASA to run on Apollo navigation computer.

Kalman Filter Modeling Assumption

- Continuous state, inputs, measurements
- **Prior over the state** is **Gaussian**
 - $p(x_0) \sim \mathcal{N}(\mu_0, \Sigma_0)$
- **Motion model**: linear with additive Gaussian disturbances
 - $x_t = A_t x_{t-1} + B_t u_t + \epsilon_t \quad \epsilon_t \sim \mathcal{N}(0, R_t)$
- Often, robotics systems are more easily described in continuous domain
 - Convert to discrete time using matrix exponential
 - Matlab contains tools to perform this conversion (c2d, d2c)
- **Measurement Model**: also linear with additive Gaussian noise
 - $y_t = C_t x_t + \delta_t \quad \delta_t \sim \mathcal{N}(0, Q_t)$
 - Can add in input dependence to match up with controls literature:
 - $y_t = C_t x_t + D_t u_t + \delta_t$

Full Model Formulation:

- **State prior**:
 - $p(x_0) \sim \mathcal{N}(\mu_0, \Sigma_0)$
- **Motion model**:
 - $x_t = A_t x_{t-1} + B_t u_t + \epsilon_t \quad \epsilon_t \sim \mathcal{N}(0, R_t)$
- **Measurement model**:
 - $y_t = C_t x_t + \delta_t \quad \delta_t \sim \mathcal{N}(0, Q_t)$

Belief is Gaussian:

- Assume **belief is Gaussian** at time t
 - $bel(x_t) \sim \mathcal{N}(\mu_t, \Sigma_t)$
 - μ_t is the best estimate of the current state at time t
 - Σ_t is the covariance, indicating the certainty in the current estimate
 - \Rightarrow the predicted belief at the next time step is also **Gaussian**
 - $\overline{bel}(x_{t+1}) \sim \mathcal{N}(\overline{\mu}_{t+1}, \overline{\Sigma}_{t+1})$
 - \Rightarrow the belief at next time step is also **Gaussian**
 - $bel(x_{t+1}) \sim \mathcal{N}(\mu_{t+1}, \Sigma_{t+1})$

Goal:

- To find belief over state as accurately as possible given all available information
 - Minimize the mean square error of the estimate (**MMSE estimator**)

$$\min E[(\mu_t - x_t)^2]$$

- Same as least square problem
- \therefore Using an unbiased estimator:

$$E[\mu_t - x_t] = 0$$

- on average, your estimate is correct!
- \therefore MMSE becomes:
 - $\min E[(\mu_t - x_t)^2] = \min E[(x_t - \mu_t)^T (x_t - \mu_t)]$
 - \equiv to minimizing **the trace of the error covariance matrix**:
 - $\Rightarrow \min \text{tr}(\Sigma_t)$

[KF Algorithm Abstract]:

- (See derivations from slides)
- At each time step t , update both sets of beliefs:

1. Prediction Update:

- Find update rule for mean, covariance of predicted belief, given input and motion model
- $\bar{\mu}_t = A_t \mu_{t-1} + B_t \mu_t \Rightarrow$ mean: noise in the motion is gone
- $\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t \Rightarrow R_t$ motion model noise variance $= \sigma_t^2$

2. Measurement Update:

- Solve MMSE optimization problem to find update rule for mean, covariance of belief given measurement model and measurement
- $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1} \Leftarrow$ **Kalman Gain**, Blending factor between prediction and measurement
- $\mu_t = \bar{\mu}_t + K_t (y_t - C_t \bar{\mu}_t)$
- $\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$

[Example]

- Temperature control
 - State is current temperature difference with outside
 - One dimensional example
 - Prior:
 - fairly certain of current temperature difference
 - $\mu_0 = 10$
 - $\Sigma_0 = 1$
 - Motion Model:
 - Decaying temperature + furnace input + disturbances (opening doors, outside effects)
 - $dt = 0.1$
 - $x_t = 0.8x_{t-1} + 3u_t + r_t$
 - $A = 0.8, B = 3$
 - $r_t \sim \mathcal{N}(0, 2)$
 - Measurement Model:
 - Directly measure the current temperature difference
 - $y(t) = x(t) + \delta_t$
 - $\delta_t \sim \mathcal{N}(0, 4)$
 - Controller design:
 - Bang bang control, based on current estimate of temperature difference
 - $u(t) = \begin{cases} 1 & \mu_t < 2 \\ 0 & \mu_t > 2 \\ u(t-1) & \text{otherwise} \end{cases}$

- **Simulation:**

```

1  for t=1:length(T)
2
3      % Select control action
4      if (t>1) u(t)=u(t-1); end
5
6      if (mu > 10)
7          u(t) = 0;
8      elseif (mu < 2);
9          u(t) = 1;
10     end
11
12     % Update state
13     e = sqrt(R)*randn(1);
14     x(t+1) = A*x(t)+ B*u(t) + e;
15
16     % Determine measurement
17     d = sqrt(Q)*randn(1);
18     y(t) = C*x(t+1) + d;

```

- **Estimation:**

```

1  % Prediction update
2  mup = A*mu + B*u(t);
3  Sp = A*Sp*A' + R;
4
5  % Measurement update
6  K = Sp*C'*inv(C*Sp*C'+Q);
7  mu = mup + K*(y(t)-C*mup);
8  S = (1-K*C)*Sp;

```

- Matrix inverse $O(n^{2.4})$, matrix multiplication $O(n^2)$
- When implementing in Matlab, `inv()` performs matrix inverse for you
- For embedded code, many libraries exist
 - Try Gnu Scientific Library, easy starting point

[Summary]

Summary:

- Follows same framework as **Bayes filter**
- Requires **linear motion** and **Gaussian disturbance**
- Requires **linear measurement** and **Gaussian noise**
- It is sufficient to update **mean** and **covariance** of **beliefs**, because they **remain Gaussian**
- **Prediction step** involves **addition of Gaussians**
- **Measurement step** seeks to **minimize** mean square error of **the estimate** (MMSE)
 - Expand out **covariance** from definition and measurement model
 - Assume form of **estimator**, **linear combination of prediction** and **measurement**
 - **Solve MMSE problem** to find **optimal linear combination**
 - **Simplify covariance** update once **gain** is found

Relation to Bayes Filter:

Relation | Problem Formulation:

- Refers to [Bayes Filter | Problem Formulation](#)
- **State prior:**
 - $bel(x_0) = p(x_0) = \mathcal{N}(\mu_0, \Sigma_0)$
- **Motion model:**
 - $p(x_t | x_{t-1}, u_t) = \mathcal{N}(A_t x_{t-1} + B_t u_t, A_t \Sigma_{t-1} A_t^T + R_t)$
- **Measurement model:**
 - $p(y_t | x_t) = \mathcal{N}(C_t x_t, Q_t)$
- **Beliefs:**
 - $bel(x_t) = \mathcal{N}(\mu_t, \Sigma_t), \quad \overline{bel}(x_t) = \mathcal{N}(\overline{\mu}_t, \overline{\Sigma}_t)$

Relation | Algorithm:

- Refers to [Bayes Filter | Algorithm](#)

1. Prediction update (Total Probability):

- Insert normal distributions:

$$\begin{aligned}\overline{bel}(x_t) &= \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1} \\ &= \eta \int e^{-1/2(x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1} (x_t - A_t x_{t-1} - B_t u_t)} e^{-1/2(x_{t-1} - \mu_{t-1})^T \Sigma_{t-1}^{-1} (x_{t-1} - \mu_{t-1})} dx_{t-1}\end{aligned}$$

- Separate out terms that depend on current state
- Manipulate remaining integral into a **Gaussian PDF** form of previous state
- Integrate over full range => to get 1
- Manipulate remaining terms and solve for Kalman prediction equations:

$$\sim \mathcal{N}(A_t \mu_{t-1} + B_t u_t, A_t \Sigma_{t-1} A_t^T + R_t)$$

2. Measurement Update (Bayes Theorem):

- Measurement update:

$$\begin{aligned}bel(x_t) &= \eta p(y_t | x_t) \overline{bel}(x_t) \\ &= \eta e^{-1/2(y_t - C_t x_t)^T Q_t^{-1} (y_t - C_t x_t)} e^{-1/2(x_t - \bar{\mu}_t)^T \bar{\Sigma}_t^{-1} (x_t - \bar{\mu}_t)}\end{aligned}$$

- Reorganize exponents and note it **remains a Gaussian**
- For any Gaussian:
 - Second derivative of exponent is inverse of covariance
 - Mean minimizes exponent
 - Set first derivative of exponent to 0 and solve
- Use this to solve for mean and covariance of belief

$$= \mathcal{N}(\bar{\mu}_t + K_t (y_t - C_t \bar{\mu}_t), (I - K_t C_t) \bar{\Sigma}_t)$$

- where K_t is the Kalman gain

Variable Summary:

- Belief mean is tradeoff between prediction and measurement
 - Kalman gain determines how to blend estimates
- If Q_t is large, inverse is small => so Kalman gain remains small
 - When measurements are high in covariance, don't trust them!
- If R_t is large, then so is predicted belief covariance, => so Kalman gain becomes large
 - When model is affected by large unknown disturbances, don't trust the predicted motion!!!

Extended Kalman Filter

- A direct generalization of the Kalman filter to nonlinear motion and measurement models
 - Relies on linearization about current estimate
 - Works well when the problem **maintains locally linear and Gaussian characteristics**
 - Computationally similar to Kalman Filter
 - Covariance can diverge when approximation is poor!

EKF | Modeling Assumption:

- **Prior** over the state is Gaussian

$$p(x_0) \sim \mathcal{N}(\mu_0, \Sigma_0)$$
- **Motion mode, nonlinear** but still with additive Gaussian disturbances

$$x_t = g(x_{t-1}, u_t) + \epsilon_t \quad \epsilon_t \sim \mathcal{N}(0, R_t)$$
- **Measurement model**, also **nonlinear** with additive Gaussian noise

$$y_t = h(x_t) + \delta_t \quad \delta_t \sim \mathcal{N}(0, Q_t)$$
- Non-linearity destroys certainty that beliefs remain Gaussian

[EKF Algorithm Abstract]

0. Linearization with First Order Taylor Series Expansion

- Only valid near point of linearization
- Motion model:

$$\begin{aligned} g(x_{t-1}, u_t) &\approx g(\mu_{t-1}, u_t) + \left. \frac{\partial}{\partial x_{t-1}} g(x_{t-1}, u_t) \right|_{x_{t-1}=\mu_{t-1}} \cdot (x_{t-1} - \mu_{t-1}) \\ &= g(\mu_{t-1}, u_t) + G_t \cdot (x_{t-1} - \mu_{t-1}) \end{aligned}$$

- Measurement model:

$$\begin{aligned} h(x_t) &\approx h(\bar{\mu}_t) + \left. \frac{\partial}{\partial x_t} h(x_t) \right|_{x_t=\bar{\mu}_t} \cdot (x_t - \bar{\mu}_t) \\ &= h(\bar{\mu}_t, u_t) + H_t \cdot (x_t - \bar{\mu}_t) \end{aligned}$$

1. Prediction Update

- $G_t = \left. \frac{\partial}{\partial x_{t-1}} g(x_{t-1}, u_t) \right|_{x_{t-1}=\mu_{t-1}}$
- $\bar{\mu}_t = g(\mu_{t-1}, u_t)$
- $\bar{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t$

2. Measurement Update

- $H_t = \left. \frac{\partial}{\partial x_t} h(x_t) \right|_{x_t=\bar{\mu}_t}$
- $K_t = \bar{\Sigma}_t H_t^T (H_t \bar{\Sigma}_t H_t^T + Q_t)^{-1}$
- $\mu_t = \bar{\mu}_t + K_t (y_t - h(\bar{\mu}_t))$
- $\Sigma_t = (I - K_t H_t) \bar{\Sigma}_t$

Sample Code:

```
1  %% Extended Kalman Filter Estimation % Prediction update
2
3  mup = Ad*mu;
4
5  Sp = Ad*S*Ad' + R;
6
7  % Measurement update
8
9  Ht = [(mup(1))/(sqrt(mup(1)^2 + mup(3)^2));
10 0; (mup(3))/(sqrt(mup(1)^2 + mup(3)^2))];
11
12 K = Sp*Ht'*inv(Ht*Sp*Ht'+Q);
13
14 mu = mup + K*(y(:,t)-sqrt(mup(1)^2 + mup(3)^2));
15
16 S = (eye(n)-K*Ht)*Sp;
17
```

Summary

- Direct extension of KF to nonlinear models
- Use Taylor series expansion to find locally linear approximations
- No longer optimal
- Most effective when covariance is low
 - Local linear approximation more likely to be accurate over range of distribution
- Covariance update may diverge
- The EKF used linearization about the predicted/previous state estimate to update the mean and covariance of the current estimate
 - Approximation of a nonlinear transformation of a Gaussian distribution by linear transformation of the mean and covariance

[Summary | Comparison Table]

	Bayes Filter	Kalman Filter	Extended Kalman Filter
Prior State	$bel(x_0) = p(x_0)$	$= \mathcal{N}(\mu_0, \Sigma_0)$	$p(x_0) \sim \mathcal{N}(\mu_0, \Sigma_0)$
Motion Model	$p(x_t x_{t-1}, u_t)$	$= \mathcal{N}(A_t x_{t-1} + B_t u_t, A_t \Sigma_{t-1} A_t^T + R_t)$	$x_t = g(x_{t-1}, u_t) + \epsilon_t, \epsilon_t \sim \mathcal{N}(0, R_t)$
Measurement Model	$p(y_t x_t)$	$= \mathcal{N}(C_t x_t, Q_t)$	$y_t = h(x_t) + \delta_t, \delta_t \sim \mathcal{N}(0, Q_t)$
Beliefs	$bel(x_t) = p(x_t y_{1:t}, u_{1:t})$	$bel(x_t) = \mathcal{N}(\mu_t, \Sigma_t), \quad \overline{bel}(x_t) = \mathcal{N}(\overline{\mu}_t, \overline{\Sigma}_t)$	Remain Gaussian
1. Prediction Update	$\overline{bel}(x_t) = \int p(x_t u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$	$= \eta \int \left[e^{-1/2(x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1}(x_t - A_t x_{t-1} - B_t u_t)} e^{-1/2(x_{t-1} - \mu_{t-1})^T \Sigma_{t-1}^{-1}(x_{t-1} - \mu_{t-1})} \right] dx_{t-1}$	
(at time t)	$\overline{bel}(x_t) = \sum p(x_t u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$	$\begin{aligned} \cdot \overline{\mu}_t &= A_t \mu_{t-1} + B_t \mu_t \\ \cdot \overline{\Sigma}_t &= A_t \Sigma_{t-1} A_t^T + R_t \end{aligned}$	$G_t = \frac{\partial}{\partial x_{t-1}} g(x_{t-1}, u_t) \Big _{x_{t-1}=\mu_{t-1}} \overline{\mu}_t = g(\mu_{t-1}, u_t) \overline{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t$
2. Measurement Update	$bel(x_t) = \eta p(y_t x_t) \overline{bel}(x_t)$	$= \eta e^{-1/2(y_t - C_t \mu_t)^T Q_t^{-1}(y_t - C_t \mu_t)} e^{-1/2(x_t - \overline{\mu}_t)^T \overline{\Sigma}_t^{-1}(x_t - \overline{\mu}_t)}$	
(at time t)	$bel(x_t) = \eta p(y_t x_t) \overline{bel}(x_t)$	$\begin{aligned} \cdot K_t &= \overline{\Sigma}_t C_t^T (C_t \overline{\Sigma}_t C_t^T + Q_t)^{-1} \\ \cdot \mu_t &= \overline{\mu}_t + K_t (y_t - C_t \overline{\mu}_t) \\ \cdot \Sigma_t &= (I - K_t C_t) \overline{\Sigma}_t \end{aligned}$	$\begin{aligned} H_t &= \frac{\partial}{\partial x_t} h(x_t) \Big _{x_t=\mu_t} K_t = \overline{\Sigma}_t H_t^T (H_t \overline{\Sigma}_t H_t^T + Q_t)^{-1} \\ \mu_t &= \overline{\mu}_t + K_t (y_t - h(\overline{\mu}_t)) \Sigma_t = (I - K_t H_t) \overline{\Sigma}_t \end{aligned}$
Summary	<p>- Bayes Filter Algorithm</p> <p>- Bayes Filter Problem Formulation</p> <p>- If state, measurements, inputs are DISCRETE, -- can directly implement Bayes Filter --- Prediction update is <i>summation</i> over discrete states --- Measurement update is <i>multiplication</i> of two vectors</p> <p>- Else, if they are CONTINUOUS -- must define model or approximation to enable computation: ---- KF (Kalman Filter) ----- Linear motion models ----- Linear measurement models ----- Additive Gaussian disturbance and noise distributions ---- EKF / UKF (Extended Kalman Filter / Unscented Kalman Filter) ----- Non-linear motion models ----- Non-linear measurement models ----- Additive Gaussian disturbance and noise distributions ---- PF (Particle Filter) ----- (Dis)continuous motion models ----- (Dis)continuous measurement models ----- General disturbance and noise distributions</p>	<p>- KF Algorithm Abstract</p> <p>- Follows same framework as Bayes filter - Requires linear motion and Gaussian disturbance - Requires linear measurement and Gaussian noise - It is sufficient to update mean and covariance of beliefs, because they remain Gaussian**Prediction step** involves addition of Gaussians - Measurement step seeks to minimize mean square error of the estimate (MMSE) -- Expand out covariance from definition and measurement model -- Assume form of estimator, linear combination of prediction and measurement - Solve MMSE problem to find optimal linear combination -- Simplify covariance update once gain is found</p>	<p>- EKF Algorithm Abstract</p> <p>- non-linear measurements and motion models - no longer optimal - Most effective when covariance is low - covariance update may diverge - Approximation of a nonlinear transformation of a Gaussian distribution by linear tranformation of the mean and covariance</p>